

# analysis

February 4, 2026

## 1 Mini Project 3: Unsupervised Discovery

### 1.1 Credit Card Customer Segmentation Analysis

**Dataset:** Credit Card Customer Segmentation

**Business Context:** A bank wants to segment credit card customers based on their 6-month usage behavior to develop targeted product offerings and marketing strategies.

**Analysis Approach:** 1. Clustering Analysis - Identify natural customer segments 2. Dimensionality Reduction - Visualize high-dimensional patterns 3. Anomaly Detection - Identify unusual customer behavior 4. Integrated Analysis - Synthesize findings for business recommendations

### 1.2 1. Data Loading and Exploration

```
[130]: # Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import umap.umap_ as umap
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
from sklearn.ensemble import IsolationForest
from sklearn.metrics import silhouette_score, silhouette_samples
import warnings
warnings.filterwarnings('ignore')

# Set random seeds for reproducibility
np.random.seed(42)

# Set plotting style
sns.set_style('whitegrid')
plt.rcParams['figure.figsize'] = (12, 6)
plt.rcParams['font.size'] = 10
```

```
[131]: # Load the dataset
df = pd.read_csv('../data/CC_GENERAL.csv')

print(f"Dataset Shape: {df.shape}")
print(f"\nFirst few rows:")
df.head()
```

Dataset Shape: (8950, 18)

First few rows:

```
[131]:   CUST_ID      BALANCE  BALANCE_FREQUENCY  PURCHASES  ONEOFF_PURCHASES \
0  C10001    40.900749        0.818182     95.40          0.00
1  C10002   3202.467416        0.909091      0.00          0.00
2  C10003   2495.148862        1.000000    773.17        773.17
3  C10004   1666.670542        0.636364   1499.00      1499.00
4  C10005    817.714335        1.000000     16.00         16.00

      INSTALLMENTS_PURCHASES  CASH_ADVANCE  PURCHASES_FREQUENCY \
0                  95.4       0.000000        0.166667
1                  0.0      6442.945483        0.000000
2                  0.0       0.000000        1.000000
3                  0.0      205.788017        0.083333
4                  0.0       0.000000        0.083333

      ONEOFF_PURCHASES_FREQUENCY  PURCHASES_INSTALLMENTS_FREQUENCY \
0            0.000000              0.083333
1            0.000000              0.000000
2            1.000000              0.000000
3            0.083333              0.000000
4            0.083333              0.000000

      CASH_ADVANCE_FREQUENCY  CASH_ADVANCE_TRX  PURCHASES_TRX  CREDIT_LIMIT \
0            0.000000             0                 2       1000.0
1            0.250000             4                 0       7000.0
2            0.000000             0                12       7500.0
3            0.083333             1                 1       7500.0
4            0.000000             0                 1       1200.0

      PAYMENTS  MINIMUM_PAYMENTS  PRC_FULL_PAYMENT  TENURE
0  201.802084      139.509787        0.000000      12
1  4103.032597     1072.340217        0.222222      12
2  622.066742      627.284787        0.000000      12
3  0.000000           Nan        0.000000      12
4  678.334763     244.791237        0.000000      12
```

```
[132]: # Dataset information
print("Dataset Info:")
df.info()

print("\nBasic Statistics:")
df.describe()
```

Dataset Info:

<class 'pandas.DataFrame'>

RangeIndex: 8950 entries, 0 to 8949

Data columns (total 18 columns):

#	Column	Non-Null Count	Dtype
0	CUST_ID	8950 non-null	str
1	BALANCE	8950 non-null	float64
2	BALANCE_FREQUENCY	8950 non-null	float64
3	PURCHASES	8950 non-null	float64
4	ONEOFF_PURCHASES	8950 non-null	float64
5	INSTALLMENTS_PURCHASES	8950 non-null	float64
6	CASH_ADVANCE	8950 non-null	float64
7	PURCHASES_FREQUENCY	8950 non-null	float64
8	ONEOFF_PURCHASES_FREQUENCY	8950 non-null	float64
9	PURCHASES_INSTALLMENTS_FREQUENCY	8950 non-null	float64
10	CASH_ADVANCE_FREQUENCY	8950 non-null	float64
11	CASH_ADVANCE_TRX	8950 non-null	int64
12	PURCHASES_TRX	8950 non-null	int64
13	CREDIT_LIMIT	8949 non-null	float64
14	PAYMENTS	8950 non-null	float64
15	MINIMUM_PAYMENTS	8637 non-null	float64
16	PRC_FULL_PAYMENT	8950 non-null	float64
17	TENURE	8950 non-null	int64

dtypes: float64(14), int64(3), str(1)

memory usage: 1.2 MB

Basic Statistics:

```
[132]:      BALANCE  BALANCE_FREQUENCY    PURCHASES  ONEOFF_PURCHASES \
count    8950.000000        8950.000000    8950.000000    8950.000000
mean     1564.474828        0.877271    1003.204834    592.437371
std      2081.531879        0.236904    2136.634782    1659.887917
min       0.000000        0.000000     0.000000     0.000000
25%     128.281915        0.888889     39.635000     0.000000
50%     873.385231        1.000000    361.280000    38.000000
75%    2054.140036        1.000000   1110.130000    577.405000
max    19043.138560        1.000000   49039.570000   40761.250000

      INSTALLMENTS_PURCHASES  CASH_ADVANCE  PURCHASES_FREQUENCY \
count    8950.000000        8950.000000    8950.000000
```

mean	411.067645	978.871112	0.490351
std	904.338115	2097.163877	0.401371
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.083333
50%	89.000000	0.000000	0.500000
75%	468.637500	1113.821139	0.916667
max	22500.000000	47137.211760	1.000000
count	8950.000000	8950.000000	8950.000000
mean	0.202458	0.364437	0.364437
std	0.298336	0.397448	0.397448
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.083333	0.166667	0.166667
75%	0.300000	0.750000	0.750000
max	1.000000	1.000000	1.000000
count	8950.000000	8950.000000	8949.000000
mean	0.135144	3.248827	14.709832
std	0.200121	6.824647	24.857649
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	1.000000
50%	0.000000	0.000000	7.000000
75%	0.222222	4.000000	17.000000
max	1.500000	123.000000	358.000000
count	8950.000000	8637.000000	8950.000000
mean	1733.143852	864.206542	0.153715
std	2895.063757	2372.446607	0.292499
min	0.000000	0.019163	0.000000
25%	383.276166	169.123707	0.000000
50%	856.901546	312.343947	0.000000
75%	1901.134317	825.485459	0.142857
max	50721.483360	76406.207520	1.000000

```
[133]: # Check for missing values
missing_values = df.isnull().sum()
print("Missing Values:")
print(missing_values[missing_values > 0])
print(f"\nTotal missing values: {df.isnull().sum().sum()}")
print(f"Percentage of missing data: {((df.isnull().sum().sum() / (df.shape[0] * df.shape[1]) * 100):.2f}%)")
```

Missing Values:

CREDIT_LIMIT	1
--------------	---

```
MINIMUM_PAYMENTS      313
dtype: int64

Total missing values: 314
Percentage of missing data: 0.19%
```

### 1.3 2. Data Preprocessing

```
[134]: # Handle missing values
# MINIMUM_PAYMENTS has missing values - we'll impute with median
# CREDIT_LIMIT also has one missing value

print("Before imputation:")
print(df[['MINIMUM_PAYMENTS', 'CREDIT_LIMIT']].isnull().sum())

# Impute missing values with median (more robust to outliers)
# IMPORTANT: Need to use loc or assign back to df for changes to take effect
df.loc[:, 'MINIMUM_PAYMENTS'] = df['MINIMUM_PAYMENTS'].fillna(df['MINIMUM_PAYMENTS'].median())
df.loc[:, 'CREDIT_LIMIT'] = df['CREDIT_LIMIT'].fillna(df['CREDIT_LIMIT'].median())

print("\nAfter imputation:")
print(df[['MINIMUM_PAYMENTS', 'CREDIT_LIMIT']].isnull().sum())

# Verify no missing values remain
print(f"\nTotal missing values after preprocessing: {df.isnull().sum().sum()}"")
```

```
Before imputation:
MINIMUM_PAYMENTS      313
CREDIT_LIMIT           1
dtype: int64
```

```
After imputation:
MINIMUM_PAYMENTS      0
CREDIT_LIMIT           0
dtype: int64
```

```
Total missing values after preprocessing: 0
```

```
[135]: # Prepare data for clustering
# Drop CUST_ID as it's just an identifier
X = df.drop('CUST_ID', axis=1)

print(f"Feature matrix shape: {X.shape}")
print(f"\nFeatures used for analysis:")
print(X.columns.tolist())
```

```
Feature matrix shape: (8950, 17)
```

```
Features used for analysis:
```

```
['BALANCE', 'BALANCE_FREQUENCY', 'PURCHASES', 'ONEOFF_PURCHASES',
'INSTALLMENTS_PURCHASES', 'CASH_ADVANCE', 'PURCHASES_FREQUENCY',
'ONEOFF_PURCHASES_FREQUENCY', 'PURCHASES_INSTALLMENTS_FREQUENCY',
'CASH_ADVANCE_FREQUENCY', 'CASH_ADVANCE_TRX', 'PURCHASES_TRX', 'CREDIT_LIMIT',
'PAYMENTS', 'MINIMUM_PAYMENTS', 'PRC_FULL_PAYMENT', 'TENURE']
```

```
[136]: # Standardize the features
# This is CRITICAL for K-means as it's distance-based
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

print("Standardization complete.")
print(f"Scaled data shape: {X_scaled.shape}")
print(f"\nMean of scaled features (should be ~0): {X_scaled.mean(axis=0)[:5]}")
print(f"Std of scaled features (should be ~1): {X_scaled.std(axis=0)[:5]}")
```

```
Standardization complete.
```

```
Scaled data shape: (8950, 17)
```

```
Mean of scaled features (should be ~0): [8.89170798e-17 2.28643919e-16
0.0000000e+00 3.17560999e-17
2.54048799e-17]
Std of scaled features (should be ~1): [1. 1. 1. 1. 1.]
```

### 1.4 3. Clustering Analysis

We'll use K-means clustering to identify customer segments. We'll determine the optimal K using:

1. **Elbow Method** - Looking for the “elbow” in the within-cluster sum of squares
2. **Silhouette Score** - Measuring how well-separated the clusters are

```
[137]: # Elbow Method - Calculate inertia for K=2 to K=10
inertias = []
silhouette_scores = []
K_range = range(2, 11)

for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(X_scaled)
    inertias.append(kmeans.inertia_)
    silhouette_scores.append(silhouette_score(X_scaled, kmeans.labels_))

print("K-means analysis complete for K=2 to K=10")
```

```
K-means analysis complete for K=2 to K=10
```

```
[138]: # Plot Elbow Curve and Silhouette Scores
fig, axes = plt.subplots(1, 2, figsize=(15, 5))


# Elbow curve
axes[0].plot(K_range, inertias, 'bo-', linewidth=2, markersize=8)
axes[0].set_xlabel('Number of Clusters (K)', fontsize=12)
axes[0].set_ylabel('Within-Cluster Sum of Squares (Inertia)', fontsize=12)
axes[0].set_title('Elbow Method for Optimal K', fontsize=14, fontweight='bold')
axes[0].grid(True, alpha=0.3)
axes[0].set_xticks(K_range)

# Silhouette scores
axes[1].plot(K_range, silhouette_scores, 'ro-', linewidth=2, markersize=8)
axes[1].set_xlabel('Number of Clusters (K)', fontsize=12)
axes[1].set_ylabel('Silhouette Score', fontsize=12)
axes[1].set_title('Silhouette Score for Different K Values', fontsize=14, fontweight='bold')
axes[1].grid(True, alpha=0.3)
axes[1].set_xticks(K_range)
axes[1].axhline(y=0.5, color='g', linestyle='--', alpha=0.5, label='Good threshold (0.5)')
axes[1].legend()

plt.tight_layout()
plt.show()

# Print the scores
print("\nDetailed Metrics:")
print("K\tInertia\tSilhouette Score")
print("-" * 50)
for k, inertia, sil_score in zip(K_range, inertias, silhouette_scores):
    print(f"{k}\t{inertia:.2f}\t{sil_score:.4f}")

```



Detailed Metrics:

K	Inertia	Silhouette Score
<hr/>		
2	127784.53	0.2100
3	111975.04	0.2510
4	99061.94	0.1977
5	91490.50	0.1931
6	84826.59	0.2029
7	79856.16	0.2077
8	74484.88	0.2217
9	69828.70	0.2260
10	66466.41	0.2204

```
[139]: import os

os.makedirs("figs", exist_ok=True)

import matplotlib.pyplot as plt


fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Elbow
axes[0].plot(K_range, inertias, 'bo-', linewidth=2)
axes[0].set_xlabel('Number of Clusters (K)')
axes[0].set_ylabel('Inertia')
axes[0].set_title('Elbow Method')
axes[0].grid(True)

# Silhouette
axes[1].plot(K_range, silhouette_scores, 'ro-', linewidth=2)
axes[1].set_xlabel('Number of Clusters (K)')
axes[1].set_ylabel('Silhouette Score')
axes[1].set_title('Silhouette Scores')
axes[1].axvline(x=3, color='black', linestyle='--', label='Selected K=3')
axes[1].legend()
axes[1].grid(True)

plt.tight_layout()

plt.savefig("../figs/elbow_silhouette.png", dpi=300, bbox_inches='tight')
plt.show()
```



### 1.4.1 Choosing Optimal K

Based on the elbow curve and silhouette scores, we need to balance:

- **Statistical quality:** Higher silhouette score is better
- **Diminishing returns:** Where the elbow occurs
- **Business interpretability:** Too many clusters are hard to action

We chose K=3 because:

1. The silhouette score is highest at 0.2510, indicating the best-defined clusters
2. The elbow curve shows significant improvement up to K=3, with diminishing returns beyond this point
3. Three segments provide clear, actionable business categories without over-segmentation
4. While the silhouette score is moderate (0.25), this is typical for customer behavioral data where segments naturally overlap

```
[140]: # Choose optimal K based on analysis
# IMPORTANT: Adjust this based on YOUR elbow curve and silhouette scores
optimal_k = 3

print(f"\n{'='*60}")
print(f"OPTIMAL K SELECTION: K = {optimal_k}")
print(f"{'='*60}")
print(f"\nJustification:")
print(f"1. Elbow Method: The elbow curve shows diminishing returns after"
      f" K={optimal_k}")
print(f"2. Silhouette Score at K={optimal_k}: {silhouette_scores[optimal_k-2]:."
      f"4f}")
print(f"3. Business Interpretability: {optimal_k} segments are actionable for"
      f" marketing")
print(f"4. Statistical Quality: Good cluster separation with manageable"
      f" complexity")
print(f"{'='*60}\n")
```

```
=====
OPTIMAL K SELECTION: K = 3
=====
```

Justification:

1. Elbow Method: The elbow curve shows diminishing returns after K=3
  2. Silhouette Score at K=3: 0.2510
  3. Business Interpretability: 3 segments are actionable for marketing
  4. Statistical Quality: Good cluster separation with manageable complexity
- 

```
[141]: # Fit final K-means model with optimal K
kmeans_final = KMeans(n_clusters=optimal_k, random_state=42, n_init=10)
cluster_labels = kmeans_final.fit_predict(X_scaled)

# Add cluster labels to original dataframe
df['Cluster'] = cluster_labels

print(f"Final clustering complete with K={optimal_k}")
print(f"\nCluster distribution:")
print(df['Cluster'].value_counts().sort_index())
print(f"\nPercentage distribution:")
print((df['Cluster'].value_counts(normalize=True).sort_index() * 100).round(2))
```

Final clustering complete with K=3

Cluster distribution:

```
Cluster
0    1275
1    6114
2    1561
Name: count, dtype: int64
```

Percentage distribution:

```
Cluster
0    14.25
1    68.31
2    17.44
Name: proportion, dtype: float64
```

#### 1.4.2 Cluster Characterization

Now let's analyze what makes each cluster unique by examining the mean values of key features:

```
[142]: # Calculate cluster centers in original scale for interpretation
cluster_centers_original = scaler.inverse_transform(kmeans_final.
    ↴cluster_centers_)
cluster_centers_df = pd.DataFrame(cluster_centers_original, columns=X.columns)
cluster_centers_df.index = [f'Cluster {i}' for i in range(optimal_k)]
```

```

print("\nCluster Centers (Original Scale):")
print(cluster_centers_df.round(2))

```

Cluster Centers (Original Scale):					
	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	\
Cluster 0	2181.00	0.98	4183.41	2661.49	
Cluster 1	807.61	0.83	495.53	246.97	
Cluster 2	4025.65	0.96	389.30	252.50	
	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY	\	
Cluster 0	1522.40	449.19	0.95		
Cluster 1	248.88	339.11	0.46		
Cluster 2	136.89	3919.44	0.23		
	ONEOFF_PURCHASES_FREQUENCY	PURCHASES_INSTALLMENTS_FREQUENCY	\		
Cluster 0	0.66		0.74		
Cluster 1	0.13		0.34		
Cluster 2	0.11		0.15		
	CASH_ADVANCE_FREQUENCY	CASH_ADVANCE_TRX	PURCHASES_TRX	\	
Cluster 0	0.06	1.51	55.66		
Cluster 1	0.07	1.24	8.47		
Cluster 2	0.45	12.55	5.64		
	CREDIT_LIMIT	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE
Cluster 0	7635.75	4072.79	1226.90	0.30	11.92
Cluster 1	3266.72	906.93	530.07	0.15	11.48
Cluster 2	6733.02	3055.52	1765.94	0.03	11.35

```

[143]: # Detailed cluster analysis - Key features
key_features = ['BALANCE', 'PURCHASES', 'CREDIT_LIMIT', 'PAYMENTS',
                 'CASH_ADVANCE', 'PURCHASES_FREQUENCY', 'TENURE']

cluster_summary = df.groupby('Cluster')[key_features].mean()
print("\nCluster Summary - Key Features:")
print(cluster_summary.round(2))

# Calculate relative to overall mean to see which clusters are above/below
# average
print("\n\nCluster Characteristics Relative to Dataset Mean:")
print("(Values >1 indicate above average, <1 indicate below average)")
cluster_relative = cluster_summary / df[key_features].mean()
print(cluster_relative.round(2))

```

Cluster Summary - Key Features:						
	BALANCE	PURCHASES	CREDIT_LIMIT	PAYMENTS	CASH_ADVANCE	\
Cluster 0	1.00	1.00	1.00	1.00	1.00	
Cluster 1	0.63	0.63	0.63	0.63	0.63	
Cluster 2	1.00	1.00	1.00	1.00	1.00	

Cluster					
0	2182.35	4187.02	7642.78	4075.53	449.75
1	807.72	496.06	3267.02	907.45	339.00
2	4023.79	389.05	6729.47	3053.94	3917.25

PURCHASES\_FREQUENCY TENURE

Cluster			
0	0.95	11.92	
1	0.46	11.48	
2	0.23	11.35	

Cluster Characteristics Relative to Dataset Mean:

(Values >1 indicate above average, <1 indicate below average)

BALANCE PURCHASES CREDIT\_LIMIT PAYMENTS CASH\_ADVANCE \


Cluster					
0	1.39	4.17	1.70	2.35	0.46
1	0.52	0.49	0.73	0.52	0.35
2	2.57	0.39	1.50	1.76	4.00

PURCHASES\_FREQUENCY TENURE

Cluster			
0	1.93	1.03	
1	0.94	1.00	
2	0.48	0.99	

[144]: # Visualize cluster characteristics with heatmap

```
plt.figure(figsize=(14, 6))
sns.heatmap(cluster_relative.T, annot=True, fmt=' .2f ', cmap='RdYlGn',
            center=1, cbar_kws={'label': 'Relative to Mean'}, linewidths=0.5)
plt.title('Cluster Characteristics Heatmap (Relative to Dataset Mean)',
          fontsize=14, fontweight='bold', pad=20)
plt.xlabel('Cluster', fontsize=12)
plt.ylabel('Feature', fontsize=12)
plt.tight_layout()
plt.show()
```



### 1.4.3 Naming Clusters with Business-Meaningful Labels

Based on the cluster characteristics, let's assign meaningful names:

```
[145]: # Name clusters based on their characteristics
# ADJUSTED based on actual cluster analysis

cluster_names = {
    0: 'Premium Shoppers',           # Cluster 0: High purchases (4.17x), high
    ↪payments, very active          # low balance, high credit limit
    1: 'Low Engagement',            # Cluster 1: Low activity across all metrics
    2: 'Cash Advance Users'        # Cluster 2: Very high cash advances (4.00x), low
    ↪low purchases
}

df['Cluster_Name'] = df['Cluster'].map(cluster_names)

print("\nCluster Names and Justifications:")
print("=" * 80)

for cluster_id in range(optimal_k):
    cluster_data = df[df['Cluster'] == cluster_id]
    print(f"\nCluster {cluster_id}: {cluster_names[cluster_id]}")
    print(f"Size: {len(cluster_data)} customers ({len(cluster_data)/len(df)*100:.1f}%)")
    print(f"\nKey Characteristics:")
    print(f" - Avg Balance: ${cluster_data['BALANCE'].mean():.2f}")
    print(f" - Avg Purchases: ${cluster_data['PURCHASES'].mean():.2f}")
    print(f" - Avg Cash Advance: ${cluster_data['CASH_ADVANCE'].mean():.2f}")
    print(f" - Avg Credit Limit: ${cluster_data['CREDIT_LIMIT'].mean():.2f}")
```

```
    print(f" - Purchase Frequency: {cluster_data['PURCHASES_FREQUENCY'].mean():.2f}")
print("—" * 80)
```

Cluster Names and Justifications:

---

Cluster 0: Premium Shoppers  
Size: 1275 customers (14.2%)

Key Characteristics:

- Avg Balance: \$2182.35
  - Avg Purchases: \$4187.02
  - Avg Cash Advance: \$449.75
  - Avg Credit Limit: \$7642.78
  - Purchase Frequency: 0.95
- 

Cluster 1: Low Engagement  
Size: 6114 customers (68.3%)

Key Characteristics:

- Avg Balance: \$807.72
  - Avg Purchases: \$496.06
  - Avg Cash Advance: \$339.00
  - Avg Credit Limit: \$3267.02
  - Purchase Frequency: 0.46
- 

Cluster 2: Cash Advance Users  
Size: 1561 customers (17.4%)

Key Characteristics:

- Avg Balance: \$4023.79
  - Avg Purchases: \$389.05
  - Avg Cash Advance: \$3917.25
  - Avg Credit Limit: \$6729.47
  - Purchase Frequency: 0.23
- 

```
[146]: # Create cluster comparison visualization
fig, axes = plt.subplots(2, 3, figsize=(18, 10))
fig.suptitle('Cluster Comparison Across Key Features', fontsize=16,
             fontweight='bold', y=1.00)
```

```

features_to_plot = ['BALANCE', 'PURCHASES', 'CASH_ADVANCE', 'CREDIT_LIMIT',
                    'PAYMENTS', 'PURCHASES_FREQUENCY']
colors = plt.cm.Set3(range(optimal_k))


for idx, feature in enumerate(features_to_plot):
    ax = axes[idx // 3, idx % 3]

    cluster_means = df.groupby('Cluster_Name')[feature].mean().
    ↪sort_values(ascending=False)
    bars = ax.bar(range(len(cluster_means)), cluster_means.values, color=colors)
    ax.set_xticks(range(len(cluster_means)))
    ax.set_xticklabels(cluster_means.index, rotation=45, ha='right', fontsize=9)
    ax.set_title(feature, fontweight='bold')
    ax.set_ylabel('Average Value')
    ax.grid(axis='y', alpha=0.3)

    # Add value labels on bars
    for bar in bars:
        height = bar.get_height()
        ax.text(bar.get_x() + bar.get_width()/2., height,
                f'{height:.0f}',
                ha='center', va='bottom', fontsize=8)

plt.tight_layout()
plt.show()

```



## 1.5 4. Dimensionality Reduction

We'll apply multiple dimensionality reduction techniques to visualize our high-dimensional data:

1. **PCA** - Linear dimensionality reduction (preserves global structure)
2. **t-SNE** - Non-linear dimensionality reduction (preserves local structure)
3. **UMAP** - Non-linear dimensionality reduction (balances local and global structure)

### 1.5.1 PCA (Principal Component Analysis)

```
[147]: # Apply PCA
pca = PCA(n_components=2, random_state=42)
X_pca = pca.fit_transform(X_scaled)

# Add PCA components to dataframe
df['PCA1'] = X_pca[:, 0]
df['PCA2'] = X_pca[:, 1]

print("PCA Analysis:")
print(f"Explained Variance Ratio: {pca.explained_variance_ratio_}")
print(f"Total Variance Explained: {pca.explained_variance_ratio_.sum():.2%}")
print(f"\nInterpretation: The first 2 principal components capture {pca.
    ↪explained_variance_ratio_.sum():.2%} ")
print(f"of the total variance in the {X.shape[1]} original features.")
```

PCA Analysis:  
Explained Variance Ratio: [0.27297671 0.2031378 ]  
Total Variance Explained: 47.61%

Interpretation: The first 2 principal components capture 47.61%  
of the total variance in the 17 original features.


```
[148]: # Visualize PCA with cluster colors
plt.figure(figsize=(12, 8))
scatter = plt.scatter(df['PCA1'], df['PCA2'], c=df['Cluster'],
                      cmap='Set3', s=50, alpha=0.6, edgecolors='black', ↪
                      linewidth=0.5)
plt.xlabel(f'First Principal Component ({pca.explained_variance_ratio_[0]:.1%} ↪
variance)', fontsize=12)
plt.ylabel(f'Second Principal Component ({pca.explained_variance_ratio_[1]:.1%} ↪
variance)', fontsize=12)
plt.title('PCA Visualization of Credit Card Customers (Colored by Cluster)',
          fontsize=14, fontweight='bold')
plt.colorbar(scatter, label='Cluster', ticks=range(optimal_k))
plt.grid(True, alpha=0.3)

# Add cluster centers
pca_centers = pca.transform(kmeans_final.cluster_centers_)
plt.scatter(pca_centers[:, 0], pca_centers[:, 1],
```

```

        c='red', s=300, alpha=0.8, edgecolors='black', linewidth=2,
        marker='X', label='Cluster Centers')
plt.legend()
plt.tight_layout()
plt.show()

```



### 1.5.2 t-SNE (t-Distributed Stochastic Neighbor Embedding)

```

[149]: # Apply t-SNE
tsne = TSNE(n_components=2, random_state=42, perplexity=30, max_iter=1000) # ↴
    ↴Changed n_iter to max_iter
X_tsne = tsne.fit_transform(X_scaled)

# Add t-SNE components to dataframe
df['tSNE1'] = X_tsne[:, 0]
df['tSNE2'] = X_tsne[:, 1]


print("t-SNE Analysis Complete")
print("Note: t-SNE is particularly good at revealing local structure and ↴
    ↴cluster separation.")

```

t-SNE Analysis Complete

Note: t-SNE is particularly good at revealing local structure and cluster separation.

```
[150]: # Visualize t-SNE with cluster colors
plt.figure(figsize=(12, 8))
scatter = plt.scatter(df['tSNE1'], df['tSNE2'], c=df['Cluster'],
                      cmap='Set3', s=50, alpha=0.6, edgecolors='black', linewidth=0.5)
plt.xlabel('t-SNE Component 1', fontsize=12)
plt.ylabel('t-SNE Component 2', fontsize=12)
plt.title('t-SNE Visualization of Credit Card Customers (Colored by Cluster)', fontsize=14, fontweight='bold')
plt.colorbar(scatter, label='Cluster', ticks=range(optimal_k))
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



### 1.5.3 UMAP (Uniform Manifold Approximation and Projection)

```
[151]: # Install UMAP if not already installed
try:
    import umap
    print("UMAP already installed")
except ImportError:
    print("Installing UMAP... This may take a minute.")
    import subprocess
    import sys

    # First upgrade pip and setuptools
    subprocess.check_call([sys.executable, "-m", "pip", "install", "--upgrade", "pip", "setuptools", "wheel"])

    # Then install umap-learn
    subprocess.check_call([sys.executable, "-m", "pip", "install", "umap-learn"])

import umap
print("UMAP installed successfully!")

# Apply UMAP
umap_reducer = umap.UMAP(n_components=2, random_state=42, n_neighbors=15, min_dist=0.1)
X_umap = umap_reducer.fit_transform(X_scaled)

# Add UMAP components to dataframe
df['UMAP1'] = X_umap[:, 0]
df['UMAP2'] = X_umap[:, 1]

print("UMAP Analysis Complete")
print("Note: UMAP balances local and global structure preservation.")
```

```
UMAP already installed
UMAP Analysis Complete
Note: UMAP balances local and global structure preservation.
```

```
[152]: import os
import numpy as np
import matplotlib.pyplot as plt

# Make folder
os.makedirs("figs", exist_ok=True)

# (Optional) nicer labels if you have them
cluster_labels = {
    0: "Premium Shoppers",
```

```

1: "Low Engagement",
2: "Cash Advance Users"
}

plt.figure(figsize=(12, 8))

scatter = plt.scatter(
    df['UMAP1'],
    df['UMAP2'],
    c=df['Cluster'],
    cmap='tab10',
    s=20,
    alpha=0.5,
    edgecolors='none'
)


plt.xlabel('UMAP Component 1', fontsize=12)
plt.ylabel('UMAP Component 2', fontsize=12)
plt.title('UMAP Projection of Customer Segments Identified by K-Means',
          fontsize=14, fontweight='bold')

plt.legend(handles=scatter.legend_elements()[0],
           labels=['Premium Shoppers', 'Low Engagement', 'Cash Advance Users'],
           title="Cluster")

plt.grid(True, alpha=0.3)
plt.tight_layout()

plt.savefig("../figs/umap_clusters.png", dpi=300, bbox_inches='tight')
plt.show()

```



#### 1.5.4 Comparison of Dimensionality Reduction Techniques

```
[153]: # Side-by-side comparison of all three techniques
fig, axes = plt.subplots(1, 3, figsize=(20, 6))
fig.suptitle('Comparison of Dimensionality Reduction Techniques',
             fontsize=16, fontweight='bold', y=1.02)

# PCA
scatter1 = axes[0].scatter(df['PCA1'], df['PCA2'], c=df['Cluster'],
                           cmap='Set3', s=30, alpha=0.6, edgecolors='black',
                           linewidth=0.3)
axes[0].set_xlabel(f'PC1 ({pca.explained_variance_ratio_[0]:.1%})')
axes[0].set_ylabel(f'PC2 ({pca.explained_variance_ratio_[1]:.1%})')
axes[0].set_title('PCA - Linear Reduction\n(Preserves Global Structure)',
                  fontweight='bold')
axes[0].grid(True, alpha=0.3)

# t-SNE
scatter2 = axes[1].scatter(df['tSNE1'], df['tSNE2'], c=df['Cluster'],
                           cmap='Set3', s=30, alpha=0.6, edgecolors='black',
                           linewidth=0.3)
axes[1].set_xlabel('t-SNE 1')
axes[1].set_ylabel('t-SNE 2')
```

```

axes[1].set_title('t-SNE - Non-linear Reduction\n(Preserves Local Structure)',  

    ↪fontweight='bold')  

axes[1].grid(True, alpha=0.3)  
  

# UMAP  

scatter3 = axes[2].scatter(df['UMAP1'], df['UMAP2'], c=df['Cluster'],  

    cmap='Set3', s=30, alpha=0.6, edgecolors='black',  

    ↪linewidth=0.3)  

axes[2].set_xlabel('UMAP 1')  

axes[2].set_ylabel('UMAP 2')  

axes[2].set_title('UMAP - Non-linear Reduction\n(Balances Local & Global)',  

    ↪fontweight='bold')  

axes[2].grid(True, alpha=0.3)  
  


# Add shared colorbar  

cbar = plt.colorbar(scatter3, ax=axes, label='Cluster', ticks=range(optimal_k))  
  

plt.tight_layout()  

plt.show()

```



### 1.5.5 Critical Analysis of Dimensionality Reduction

#### Key Observations:

1. **PCA Analysis:**
  - Shows the linear relationships in the data
  - The first 2 components explain **47.6%** of variance
  - Clusters show **moderate overlap**, with Premium Shoppers (Cluster 0) slightly separated
  - Best for: Understanding major axes of variation and quantifying explained variance
2. **t-SNE Analysis:**
  - Reveals local cluster structure very clearly
  - Shows **distinct separation between all three clusters**
  - Premium Shoppers and Cash Advance Users form tighter groups
  - Low Engagement cluster is more diffuse (expected for inactive customers)

- Warning: Distances between clusters less meaningful
- Best for: Visualizing distinct groups and cluster boundaries

### 3. UMAP Analysis:

- Balances local and global structure
- Shows **similar patterns to t-SNE but with better global structure preservation**
- Reveals that Low Engagement cluster has internal substructure
- More interpretable distances than t-SNE
- Best for: General-purpose visualization and business presentations

**Recommendation for Business Presentation:** - For executive presentation, I would use **UMAP** because it provides clear cluster separation while maintaining interpretable global structure. The visualization is intuitive without being overly complex. - The data appears to have **non-linear structure** based on the fact that t-SNE and UMAP reveal clearer patterns than PCA (which only captures 47.6% variance) - Cluster separation is **moderately well-separated**, indicating genuine behavioral differences while acknowledging that customers exist on a continuum rather than in discrete categories. - For business executives, I recommend UMAP because it provides clear cluster separation while maintaining interpretable distances, making it easy to explain the three customer segments without technical jargon.

## 1.6 5. Anomaly Detection

We'll use Isolation Forest to detect unusual customer behaviors that don't fit typical patterns.

```
[154]: # Test different contamination values
contamination_values = [0.01, 0.05, 0.10]

print("Testing different contamination values:")
print("=" * 60)

for contamination in contamination_values:
    iso_forest = IsolationForest(contamination=contamination, random_state=42)
    anomaly_labels = iso_forest.fit_predict(X_scaled)
    n_anomalies = (anomaly_labels == -1).sum()

    print(f"\nContamination = {contamination}:")
    print(f"  Number of anomalies: {n_anomalies} ({n_anomalies/len(df)*100:.2f}%)")
    print(f"  Cluster distribution of anomalies:")
    anomaly_clusters = df[anomaly_labels == -1]['Cluster'].value_counts()
    .sort_index()
    for cluster, count in anomaly_clusters.items():
        print(f"    Cluster {cluster} ({cluster_names[cluster]}): {count} anomalies")
```

Testing different contamination values:

```
=====
Contamination = 0.01:
Number of anomalies: 90 (1.01%)
```

```
Cluster distribution of anomalies:  
    Cluster 0 (Premium Shoppers): 72 anomalies  
    Cluster 2 (Cash Advance Users): 18 anomalies  
  
Contamination = 0.05:  
    Number of anomalies: 448 (5.01%)  
    Cluster distribution of anomalies:  
        Cluster 0 (Premium Shoppers): 282 anomalies  
        Cluster 1 (Low Engagement): 8 anomalies  
        Cluster 2 (Cash Advance Users): 158 anomalies  
  
Contamination = 0.1:  
    Number of anomalies: 895 (10.00%)  
    Cluster distribution of anomalies:  
        Cluster 0 (Premium Shoppers): 488 anomalies  
        Cluster 1 (Low Engagement): 92 anomalies  
        Cluster 2 (Cash Advance Users): 315 anomalies
```

FINAL CONTAMINATION CHOICE: 0.05

---

Justification:

1. Domain knowledge: ~5% unusual behavior is reasonable for credit card data
  2. Balance: Not too conservative (missing real anomalies) or liberal (false positives)
  3. Actionability: Manageable number of cases to investigate
- 

Anomaly detection complete:

Total anomalies detected: 448 (5.01%)

```
[156]: # Analyze anomalies in detail
anomalies = df[df['Is_Anomaly']]

print("\nAnomaly Analysis:")
print("==" * 80)
print(f"\nTotal Anomalies: {len(anomalies)}")
print(f"\nCluster Distribution of Anomalies:")
print(anomalies['Cluster_Name'].value_counts())
print(f"\nPercentage of each cluster that are anomalies:")
for cluster_id in range(optimal_k):
    cluster_total = len(df[df['Cluster'] == cluster_id])
    cluster_anomalies = len(anomalies[anomalies['Cluster'] == cluster_id])
    pct = cluster_anomalies / cluster_total * 100 if cluster_total > 0 else 0
    print(f"  {cluster_names[cluster_id]}: {pct:.1f}% ({cluster_anomalies}/
        {cluster_total})")
```

Anomaly Analysis:

---

Total Anomalies: 448

Cluster Distribution of Anomalies:

Cluster_Name	count
Premium Shoppers	282
Cash Advance Users	158
Low Engagement	8

Name: count, dtype: int64

Percentage of each cluster that are anomalies:

Cluster	Percentage
Premium Shoppers	22.1% (282/1275)
Low Engagement	0.1% (8/6114)
Cash Advance Users	10.1% (158/1561)

```
[157]: # Compare anomalies vs normal customers
print("\nComparison: Anomalies vs Normal Customers")
print("=" * 80)

comparison_features = ['BALANCE', 'PURCHASES', 'CASH_ADVANCE', 'CREDIT_LIMIT',
                      'PAYMENTS', 'PURCHASES_FREQUENCY']

comparison_df = pd.DataFrame({
    'Normal': df[~df['Is_Anomaly']][comparison_features].mean(),
    'Anomalies': anomalies[comparison_features].mean(),
    'Ratio': anomalies[comparison_features].mean() / df[~df['Is_Anomaly']][comparison_features].mean()
})

print(comparison_df.round(2))
```

Comparison: Anomalies vs Normal Customers

	Normal	Anomalies	Ratio
BALANCE	1403.71	4615.50	3.29
PURCHASES	734.64	6099.86	8.30
CASH_ADVANCE	826.93	3862.32	4.67
CREDIT_LIMIT	4198.96	10098.82	2.41
PAYMENTS	1357.49	8862.14	6.53
PURCHASES_FREQUENCY	0.47	0.80	1.68

```
[158]: # Examine specific anomaly examples
print("\nDetailed Examples of Anomalous Customers:")
print("=" * 80)

# Get top 5 most anomalous (most negative scores)
top_anomalies = anomalies.nsmallest(5, 'Anomaly_Score')

for idx, (_, row) in enumerate(top_anomalies.iterrows(), 1):
    print(f"\nAnomaly Example {idx}:")

    print(f"  Customer ID: {row['CUST_ID']}")

    print(f"  Cluster: {row['Cluster_Name']}")

    print(f"  Anomaly Score: {row['Anomaly_Score']:.4f}")

    print(f"  Balance: ${row['BALANCE']:.2f}")

    print(f"  Purchases: ${row['PURCHASES']:.2f}")

    print(f"  Cash Advance: ${row['CASH_ADVANCE']:.2f}")

    print(f"  Credit Limit: ${row['CREDIT_LIMIT']:.2f}")

    print(f"  Payments: ${row['PAYMENTS']:.2f}")

    print("-" * 80)
```

Detailed Examples of Anomalous Customers:

---

Anomaly Example 1:

Customer ID: C10144  
Cluster: Premium Shoppers  
Anomaly Score: -0.7519  
Balance: \$19043.14  
Purchases: \$22009.92  
Cash Advance: \$0.00  
Credit Limit: \$18000.00  
Payments: \$23018.58

---

Anomaly Example 2:

Customer ID: C10523  
Cluster: Premium Shoppers  
Anomaly Score: -0.7461  
Balance: \$13479.29  
Purchases: \$41050.40  
Cash Advance: \$0.00  
Credit Limit: \$17000.00  
Payments: \$36066.75

---

Anomaly Example 3:

Customer ID: C10159  
Cluster: Premium Shoppers  
Anomaly Score: -0.7449  
Balance: \$13673.08  
Purchases: \$9792.23  
Cash Advance: \$2444.45  
Credit Limit: \$20000.00  
Payments: \$11717.31

---

Anomaly Example 4:

Customer ID: C10574  
Cluster: Premium Shoppers  
Anomaly Score: -0.7431  
Balance: \$11547.52  
Purchases: \$49039.57  
Cash Advance: \$558.17  
Credit Limit: \$22500.00  
Payments: \$46930.60

---

Anomaly Example 5:

Customer ID: C14640

```

Cluster: Premium Shoppers
Anomaly Score: -0.7429
Balance: $6956.38
Purchases: $11500.94
Cash Advance: $15133.53
Credit Limit: $14000.00
Payments: $20122.01
-----
```

### 1.6.1 Categorizing Anomalies

Let's categorize anomalies into meaningful business types:

```
[159]: # Categorize anomalies based on their characteristics
def categorize_anomaly(row):
    # High-value customers (VIPs)
    if row['CREDIT_LIMIT'] > df['CREDIT_LIMIT'].quantile(0.95) and \
       row['PURCHASES'] > df['PURCHASES'].quantile(0.90):
        return 'VIP High Spender'

    # High cash advance users
    elif row['CASH_ADVANCE'] > df['CASH_ADVANCE'].quantile(0.90):
        return 'Heavy Cash Advance User'

    # Unusual payment patterns
    elif row['BALANCE'] > df['BALANCE'].quantile(0.90) and row['PAYMENTS'] < \
         df['PAYMENTS'].quantile(0.10):
        return 'High Balance Low Payment'

    # Low engagement but high limit
    elif row['CREDIT_LIMIT'] > df['CREDIT_LIMIT'].quantile(0.75) and \
         row['PURCHASES_FREQUENCY'] < 0.1:
        return 'Inactive High Limit'

    else:
        return 'Other Unusual Pattern'

anomalies['Anomaly_Type'] = anomalies.apply(categorize_anomaly, axis=1)

print("\nAnomaly Type Distribution:")
print(anomalies['Anomaly_Type'].value_counts())
print(f"\nPercentage breakdown:")
print((anomalies['Anomaly_Type'].value_counts(normalize=True) * 100).round(1))
```

```

Anomaly Type Distribution:
Anomaly_Type
Other Unusual Pattern      189
```

```

Heavy Cash Advance User      163
VIP High Spender           91
Inactive High Limit          3
High Balance Low Payment      2
Name: count, dtype: int64

```

Percentage breakdown:

```

Anomaly_Type
Other Unusual Pattern      42.2
Heavy Cash Advance User      36.4
VIP High Spender            20.3
Inactive High Limit           0.7
High Balance Low Payment       0.4
Name: proportion, dtype: float64

```

```

[160]: # Business recommendations for each anomaly type
print("\nBusiness Actions for Each Anomaly Type:")
print("==" * 80)

anomaly_actions = {
    'VIP High Spender': 'Assign dedicated account manager, offer premium rewards, exclusive benefits',
    'Heavy Cash Advance User': 'Monitor for financial distress, offer debt consolidation, budgeting tools',
    'High Balance Low Payment': 'Flag for credit risk review, send payment reminders, offer payment plans',
    'Inactive High Limit': 'Re-engagement campaign, targeted offers, consider limit adjustment',
    'Other Unusual Pattern': 'Manual review for data quality or fraud detection'
}

for anomaly_type in anomalies['Anomaly_Type'].unique():
    count = (anomalies['Anomaly_Type'] == anomaly_type).sum()
    print(f"\n{anomaly_type} ({count} customers):")
    print(f"  Action: {anomaly_actions.get(anomaly_type, 'Review individually')}")

```

Business Actions for Each Anomaly Type:

---

VIP High Spender (91 customers):

Action: Assign dedicated account manager, offer premium rewards, exclusive benefits

Other Unusual Pattern (189 customers):

Action: Manual review for data quality or fraud detection

Heavy Cash Advance User (163 customers):  
Action: Monitor for financial distress, offer debt consolidation, budgeting tools

High Balance Low Payment (2 customers):  
Action: Flag for credit risk review, send payment reminders, offer payment plans

Inactive High Limit (3 customers):  
Action: Re-engagement campaign, targeted offers, consider limit adjustment

### 1.6.2 Visualizing Anomalies on Dimensionality Reduction Plots

```
[161]: # Visualize anomalies on all three dimensionality reduction plots
fig, axes = plt.subplots(1, 3, figsize=(20, 6))
fig.suptitle('Anomaly Detection Across Dimensionality Reduction Techniques',
             fontsize=16, fontweight='bold', y=1.02)

# PCA with anomalies
axes[0].scatter(df[~df['Is_Anomaly']]['PCA1'], df[~df['Is_Anomaly']]['PCA2'],
                c=df[~df['Is_Anomaly']]['Cluster'], cmap='Set3', s=30, alpha=0.
                ↪4, label='Normal')
axes[0].scatter(anomalies['PCA1'], anomalies['PCA2'],
                c='red', s=100, alpha=0.8, marker='X', edgecolors='black', ↪
                linewidth=1, label='Anomaly')
axes[0].set_xlabel('PCA 1')
axes[0].set_ylabel('PCA 2')
axes[0].set_title('PCA with Anomalies', fontweight='bold')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# t-SNE with anomalies
axes[1].scatter(df[~df['Is_Anomaly']]['tSNE1'], df[~df['Is_Anomaly']]['tSNE2'],
                c=df[~df['Is_Anomaly']]['Cluster'], cmap='Set3', s=30, alpha=0.
                ↪4, label='Normal')
axes[1].scatter(anomalies['tSNE1'], anomalies['tSNE2'],
                c='red', s=100, alpha=0.8, marker='X', edgecolors='black', ↪
                linewidth=1, label='Anomaly')
axes[1].set_xlabel('t-SNE 1')
axes[1].set_ylabel('t-SNE 2')
axes[1].set_title('t-SNE with Anomalies', fontweight='bold')
axes[1].legend()
axes[1].grid(True, alpha=0.3)


# UMAP with anomalies
axes[2].scatter(df[~df['Is_Anomaly']]['UMAP1'], df[~df['Is_Anomaly']]['UMAP2'],
```

```

c=df[~df['Is_Anomaly']]['Cluster'], cmap='Set3', s=30, alpha=0.
    ↵4, label='Normal')
axes[2].scatter(anomalies['UMAP1'], anomalies['UMAP2'],
                 c='red', s=100, alpha=0.8, marker='X', edgecolors='black', u
    ↵linewidth=1, label='Anomaly')
axes[2].set_xlabel('UMAP 1')
axes[2].set_ylabel('UMAP 2')
axes[2].set_title('UMAP with Anomalies', fontweight='bold')
axes[2].legend()
axes[2].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



```

[162]: import os
os.makedirs("figs", exist_ok=True)

plt.figure(figsize=(8,6))

# Normal points
normal = df['Is_Anomaly'] == False
plt.scatter(
    df.loc[normal,'UMAP1'],
    df.loc[normal,'UMAP2'],
    c=df.loc[normal,'Cluster'],
    cmap='Set2',
    s=15,
    alpha=0.6,
    label="Normal"
)

# Anomalies
anomaly = df['Is_Anomaly'] == True
plt.scatter(

```


```

df.loc[anomaly, 'UMAP1'],
df.loc[anomaly, 'UMAP2'],
c='red',
marker='x',
s=60,
label="Anomaly"
)

plt.title("UMAP Projection with Anomalies Highlighted", fontsize=14,
          weight='bold')
plt.xlabel("UMAP 1")
plt.ylabel("UMAP 2")
plt.legend()
plt.grid(alpha=0.3)

plt.tight_layout()
plt.savefig("../figs/umap_anomalies.png", dpi=300, bbox_inches='tight')
plt.show()

```



## 1.7 6. Integrated Analysis

This section synthesizes findings from all three techniques to provide comprehensive business insights.

```
[163]: # Create comprehensive summary
print("\n" + "=" * 80)
print("INTEGRATED ANALYSIS SUMMARY")
print("=" * 80)

print(f"\n1. CLUSTERING FINDINGS:")
print(f" - Identified {optimal_k} distinct customer segments")
print(f" - Segments range from {df.groupby('Cluster').size().min()} to {df.
    ↪groupby('Cluster').size().max()} customers")
print(f" - Silhouette score of {silhouette_scores[optimal_k-2]:.3f} indicates
    ↪{('good' if silhouette_scores[optimal_k-2] > 0.5 else 'moderate')} cluster
    ↪separation")

print(f"\n2. DIMENSIONALITY REDUCTION INSIGHTS:")
print(f" - PCA: First 2 components explain {pca.explained_variance_ratio_.
    ↪sum():.1%} of variance")
print(f" - t-SNE: Reveals clear local cluster structure")
print(f" - UMAP: Balances global and local patterns")
print(f" - Conclusion: Data has [linear/non-linear] structure with
    ↪[well-separated/overlapping] clusters")

print(f"\n3. ANOMALY DETECTION RESULTS:")
print(f" - Detected {len(anomalies)} anomalies ({len(anomalies)/len(df)*100:.
    ↪1f}% of customers}")
print(f" - Anomaly distribution across clusters:")
for cluster_id in range(optimal_k):
    cluster_anomalies = len(anomalies[anomalies['Cluster'] == cluster_id])
    print(f"   {cluster_names[cluster_id]}: {cluster_anomalies} anomalies")

print(f"\n4. KEY INTEGRATION INSIGHTS:")
print(f" - Dimensionality reduction confirms clustering validity: All three
    ↪methods")
print(f" - (PCA, t-SNE, UMAP) show consistent separation between clusters")
print(f" - Anomalies are heavily concentrated in Premium Shoppers (22.1%)")
print(f" - Cash Advance Users (10.1%), while Low Engagement has minimal
    ↪anomalies (0.1%)")
print(f" - This pattern makes business sense: extreme behaviors (very high
    ↪spending")
print(f" - or high cash dependence) are more likely to be anomalous")
print(f" - The moderate silhouette score (0.251) combined with clear visual")
```

```

print(f"      separation in t-SNE/UMAP indicates that while customers exist on a")
print(f"      continuum, the three segments represent meaningful behavioral categories")
print(f" - Premium Shoppers' high anomaly rate (22.1%) suggests this cluster")
print(f" contains both typical high-spenders and exceptional VIPs needing")
print(f" differentiated service")

print("\n" + "=" * 80)

```

---

## INTEGRATED ANALYSIS SUMMARY

---

### 1. CLUSTERING FINDINGS:

- Identified 3 distinct customer segments
- Segments range from 1275 to 6114 customers
- Silhouette score of 0.251 indicates moderate cluster separation

### 2. DIMENSIONALITY REDUCTION INSIGHTS:

- PCA: First 2 components explain 47.6% of variance
- t-SNE: Reveals clear local cluster structure
- UMAP: Balances global and local patterns
- Conclusion: Data has [linear/non-linear] structure with [well-separated/overlapping] clusters

### 3. ANOMALY DETECTION RESULTS:

- Detected 448 anomalies (5.0% of customers)
- Anomaly distribution across clusters:
  - Premium Shoppers: 282 anomalies
  - Low Engagement: 8 anomalies
  - Cash Advance Users: 158 anomalies

### 4. KEY INTEGRATION INSIGHTS:

- Dimensionality reduction confirms clustering validity: All three methods (PCA, t-SNE, UMAP) show consistent separation between clusters
- Anomalies are heavily concentrated in Premium Shoppers (22.1%) and Cash Advance Users (10.1%), while Low Engagement has minimal anomalies (0.1%)
- This pattern makes business sense: extreme behaviors (very high spending or high cash dependence) are more likely to be anomalous
- The moderate silhouette score (0.251) combined with clear visual separation in t-SNE/UMAP indicates that while customers exist on a continuum, the three segments represent meaningful behavioral categories
- Premium Shoppers' high anomaly rate (22.1%) suggests this cluster contains both typical high-spenders and exceptional VIPs needing

differentiated service

```
=====
```

[164]: # Create final visualization

```
fig = plt.figure(figsize=(20, 12))
gs = fig.add_gridspec(3, 3, hspace=0.3, wspace=0.3)

# Title
fig.suptitle('Unsupervised Learning Analysis: Credit Card Customer Segmentation',
             fontsize=18, fontweight='bold', y=0.98)

# Row 1: Clustering metrics
ax1 = fig.add_subplot(gs[0, 0])
ax1.plot(K_range, inertias, 'bo-', linewidth=2, markersize=8)
ax1.set_xlabel('K')
ax1.set_ylabel('Inertia')
ax1.set_title('Elbow Method')
ax1.grid(True, alpha=0.3)

ax2 = fig.add_subplot(gs[0, 1])
ax2.plot(K_range, silhouette_scores, 'ro-', linewidth=2, markersize=8)
ax2.set_xlabel('K')
ax2.set_ylabel('Silhouette Score')
ax2.set_title('Silhouette Scores')
ax2.grid(True, alpha=0.3)

ax3 = fig.add_subplot(gs[0, 2])
cluster_sizes = df['Cluster_Name'].value_counts()
ax3.bar(range(len(cluster_sizes)), cluster_sizes.values, color=plt.cm.Set3(range(optimal_k)))
ax3.set_xticks(range(len(cluster_sizes)))
ax3.set_xticklabels(cluster_sizes.index, rotation=45, ha='right')
ax3.set_ylabel('Number of Customers')
ax3.set_title('Cluster Sizes')
ax3.grid(axis='y', alpha=0.3)

# Row 2: Dimensionality reduction
ax4 = fig.add_subplot(gs[1, 0])
scatter4 = ax4.scatter(df['PCA1'], df['PCA2'], c=df['Cluster'], cmap='Set3', s=20, alpha=0.6)
ax4.set_xlabel('PCA 1')
ax4.set_ylabel('PCA 2')
ax4.set_title('PCA Visualization')
ax4.grid(True, alpha=0.3)
```

```

ax5 = fig.add_subplot(gs[1, 1])
ax5.scatter(df['tSNE1'], df['tSNE2'], c=df['Cluster'], cmap='Set3', s=20, alpha=0.6)
ax5.set_xlabel('t-SNE 1')
ax5.set_ylabel('t-SNE 2')
ax5.set_title('t-SNE Visualization')
ax5.grid(True, alpha=0.3)

ax6 = fig.add_subplot(gs[1, 2])
ax6.scatter(df['UMAP1'], df['UMAP2'], c=df['Cluster'], cmap='Set3', s=20, alpha=0.6)
ax6.set_xlabel('UMAP 1')
ax6.set_ylabel('UMAP 2')
ax6.set_title('UMAP Visualization')
ax6.grid(True, alpha=0.3)

# Row 3: Anomaly analysis
ax7 = fig.add_subplot(gs[2, 0])
anomaly_by_cluster = df.groupby('Cluster_Name')['Is_Anomaly'].sum()
ax7.bar(range(len(anomaly_by_cluster)), anomaly_by_cluster.values, color='coral')
ax7.set_xticks(range(len(anomaly_by_cluster)))
ax7.set_xticklabels(anomaly_by_cluster.index, rotation=45, ha='right')
ax7.set_ylabel('Number of Anomalies')
ax7.set_title('Anomalies by Cluster')
ax7.grid(axis='y', alpha=0.3)

ax8 = fig.add_subplot(gs[2, 1])
if 'Anomaly_Type' in anomalies.columns:
    anomaly_types = anomalies['Anomaly_Type'].value_counts()
    ax8.barch(len(anomaly_types)), anomaly_types.values, color='lightcoral')
    ax8.set_yticks(range(len(anomaly_types)))
    ax8.set_yticklabels(anomaly_types.index)
    ax8.set_xlabel('Count')
    ax8.set_title('Anomaly Types')
    ax8.grid(axis='x', alpha=0.3)

ax9 = fig.add_subplot(gs[2, 2])
ax9.scatter(~df['Is_Anomaly']][['UMAP1'], df[~df['Is_Anomaly']]][['UMAP2'],
            c=df[~df['Is_Anomaly']]][['Cluster'], cmap='Set3', s=20, alpha=0.3, label='Normal')
ax9.scatter(anomalies['UMAP1'], anomalies['UMAP2'],
            c='red', s=50, alpha=0.8, marker='X', edgecolors='black',
            linewidth=0.5, label='Anomaly')
ax9.set_xlabel('UMAP 1')

```


```

ax9.set_ylabel('UMAP 2')
ax9.set_title('Anomalies in UMAP Space')
ax9.legend()
ax9.grid(True, alpha=0.3)

plt.savefig('../figs/analysis.png', dpi=300, bbox_inches='tight')
plt.show()

print("Visualization saved as 'analysis.png'")

```



Visualization saved as 'analysis.png'

## 1.8 7. Business Recommendations

Based on the integrated analysis, here are actionable business recommendations:

```

[165]: print("\n" + "=" * 80)
print("BUSINESS RECOMMENDATIONS")
print("=" * 80)

for cluster_id in range(optimal_k):
    cluster_data = df[df['Cluster'] == cluster_id]

```

```

cluster_name = cluster_names[cluster_id]
size = len(cluster_data)
pct = size / len(df) * 100

print(f"\n{cluster_name} (Cluster {cluster_id}):")
print(f"Size: {size} customers ({pct:.1f}%)")
print(f"\nCharacteristics:")
print(f" - Avg Balance: ${cluster_data['BALANCE'].mean():.2f}")
print(f" - Avg Purchases: ${cluster_data['PURCHASES'].mean():.2f}")
print(f" - Avg Cash Advance: ${cluster_data['CASH_ADVANCE'].mean():.2f}")
print(f" - Purchase Frequency: {cluster_data['PURCHASES_FREQUENCY'].mean():.2f}")

print(f"\nRecommended Actions:")
# Add specific recommendations based on cluster characteristics
if cluster_data['PURCHASES'].mean() > df['PURCHASES'].mean() * 1.5:
    print(f"   Offer premium rewards program")
    print(f"   Provide exclusive shopping benefits")
if cluster_data['CASH_ADVANCE'].mean() > df['CASH_ADVANCE'].mean() * 1.5:
    print(f"   Offer debt consolidation products")
    print(f"   Provide financial wellness resources")
if cluster_data['PURCHASES_FREQUENCY'].mean() < 0.3:
    print(f"   Launch re-engagement campaign")
    print(f"   Send targeted promotional offers")
if cluster_data['CREDIT_LIMIT'].mean() > df['CREDIT_LIMIT'].mean() * 1.2:
    print(f"   Offer investment products")
    print(f"   Provide concierge services")

print("-" * 80)

print(f"\nANOMALY-SPECIFIC RECOMMENDATIONS:")
print(f"   VIP customers: Assign dedicated account managers")
print(f"   High cash advance users: Monitor for financial distress")
print(f"   Inactive high-limit customers: Launch win-back campaigns")
print(f"   Unusual patterns: Flag for fraud detection review")

print("\n" + "=" * 80)

```

---

## BUSINESS RECOMMENDATIONS

---

Premium Shoppers (Cluster 0):  
 Size: 1275 customers (14.2%)

Characteristics:  
 - Avg Balance: \$2182.35

- Avg Purchases: \$4187.02
- Avg Cash Advance: \$449.75
- Purchase Frequency: 0.95

Recommended Actions:

- Offer premium rewards program
  - Provide exclusive shopping benefits
  - Offer investment products
  - Provide concierge services
- 

Low Engagement (Cluster 1):

Size: 6114 customers (68.3%)

Characteristics:

- Avg Balance: \$807.72
- Avg Purchases: \$496.06
- Avg Cash Advance: \$339.00
- Purchase Frequency: 0.46

Recommended Actions:

---

Cash Advance Users (Cluster 2):

Size: 1561 customers (17.4%)

Characteristics:

- Avg Balance: \$4023.79
- Avg Purchases: \$389.05
- Avg Cash Advance: \$3917.25
- Purchase Frequency: 0.23

Recommended Actions:

- Offer debt consolidation products
  - Provide financial wellness resources
  - Launch re-engagement campaign
  - Send targeted promotional offers
  - Offer investment products
  - Provide concierge services
- 

#### ANOMALY-SPECIFIC RECOMMENDATIONS:

- VIP customers: Assign dedicated account managers
  - High cash advance users: Monitor for financial distress
  - Inactive high-limit customers: Launch win-back campaigns
  - Unusual patterns: Flag for fraud detection review
-

```
[166]: # Executive Summary
print("\n" + "="*80)
print("EXECUTIVE SUMMARY: CREDIT CARD CUSTOMER SEGMENTATION")
print("="*80)

print("\n ANALYSIS OVERVIEW:")
print(f"    Dataset: 8,950 credit card customers")
print(f"    Time Period: 6 months of transaction history")
print(f"    Features Analyzed: 17 behavioral metrics")

print("\n KEY FINDINGS:")
print(f"\n1. THREE DISTINCT CUSTOMER SEGMENTS IDENTIFIED:")
print(f"    • Premium Shoppers (14.2%): High-value, active users")
print(f"        - Average purchases: $4,187 (4.17× dataset mean)")
print(f"        - Average credit limit: $7,643")
print(f"        - Purchase frequency: 95%")
print(f"        → Revenue drivers, deserve premium service")

print(f"\n    • Low Engagement (68.3%): Underutilized accounts")
print(f"        - Average purchases: $496 (0.49× dataset mean)")
print(f"        - Purchase frequency: 46%")
print(f"        → Largest segment, biggest growth opportunity")

print(f"\n    • Cash Advance Users (17.4%): High cash dependency")
print(f"        - Average cash advance: $3,917 (4.00× dataset mean)")
print(f"        - Average balance: $4,024")
print(f"        → Credit risk concern, need intervention")

print(f"\n2. ANOMALY DETECTION RESULTS:")
print(f"    • 448 anomalous customers identified (5.0%)")
print(f"    • 91 VIP High Spenders → Dedicated account management")
print(f"    • 163 Heavy Cash Advance Users → Financial wellness programs")
print(f"    • 189 Unusual Patterns → Fraud investigation")

print(f"\n3. DATA STRUCTURE:")
print(f"    • Non-linear relationships (PCA: 47.6% variance)")
print(f"    • Moderate cluster separation (Silhouette: 0.251)")
print(f"    • Clear visual separation in t-SNE/UMAP confirms valid segments")

print("\n RECOMMENDED ACTIONS:")
print(f"    1. IMMEDIATE (Month 1):")
print(f"        - Launch VIP program for 91 identified high-spenders")
print(f"        - Implement credit risk monitoring for 163 heavy cash users")
print(f"        - Begin A/B testing re-engagement campaigns on 6,114 low-activity")
print(f"        ↴customers")

print(f"\n    2. SHORT-TERM (Months 2-6):")
```

```

print(f"      - Develop segment-specific marketing strategies")
print(f"      - Create automated cluster assignment for new customers")
print(f"      - Track monthly cluster migration patterns")

print(f"\n  3. LONG-TERM (6+ months):")
print(f"      - Build predictive models for churn and upgrade likelihood")
print(f"      - Integrate demographic data to enrich segmentation")
print(f"      - Implement real-time anomaly detection system")

print("\n EXPECTED BUSINESS IMPACT:")
print(f"  • Revenue Growth: Target 15-20% increase from re-engaged Low\u2192Engagement customers")
print(f"  • Customer Retention: Reduce churn by 10% through targeted\u2192interventions")
print(f"  • Risk Mitigation: Early identification of 163 at-risk accounts")
print(f"  • VIP Satisfaction: Enhanced service for top 91 customers")

print("\n" + "="*80)

```

=====

## EXECUTIVE SUMMARY: CREDIT CARD CUSTOMER SEGMENTATION

=====

### ANALYSIS OVERVIEW:

Dataset: 8,950 credit card customers  
 Time Period: 6 months of transaction history  
 Features Analyzed: 17 behavioral metrics

### KEY FINDINGS:

#### 1. THREE DISTINCT CUSTOMER SEGMENTS IDENTIFIED:

- Premium Shoppers (14.2%): High-value, active users
  - Average purchases: \$4,187 (4.17× dataset mean)
  - Average credit limit: \$7,643
  - Purchase frequency: 95%
  - Revenue drivers, deserve premium service
- Low Engagement (68.3%): Underutilized accounts
  - Average purchases: \$496 (0.49× dataset mean)
  - Purchase frequency: 46%
  - Largest segment, biggest growth opportunity
- Cash Advance Users (17.4%): High cash dependency
  - Average cash advance: \$3,917 (4.00× dataset mean)
  - Average balance: \$4,024
  - Credit risk concern, need intervention

2. ANOMALY DETECTION RESULTS:

- 448 anomalous customers identified (5.0%)
- 91 VIP High Spenders → Dedicated account management
- 163 Heavy Cash Advance Users → Financial wellness programs
- 189 Unusual Patterns → Fraud investigation

3. DATA STRUCTURE:

- Non-linear relationships (PCA: 47.6% variance)
- Moderate cluster separation (Silhouette: 0.251)
- Clear visual separation in t-SNE/UMAP confirms valid segments

RECOMMENDED ACTIONS:

1. IMMEDIATE (Month 1):

- Launch VIP program for 91 identified high-spenders
- Implement credit risk monitoring for 163 heavy cash users
- Begin A/B testing re-engagement campaigns on 6,114 low-activity customers

2. SHORT-TERM (Months 2-6):

- Develop segment-specific marketing strategies
- Create automated cluster assignment for new customers
- Track monthly cluster migration patterns

3. LONG-TERM (6+ months):

- Build predictive models for churn and upgrade likelihood
- Integrate demographic data to enrich segmentation
- Implement real-time anomaly detection system

EXPECTED BUSINESS IMPACT:

- Revenue Growth: Target 15-20% increase from re-engaged Low Engagement customers
  - Customer Retention: Reduce churn by 10% through targeted interventions
  - Risk Mitigation: Early identification of 163 at-risk accounts
  - VIP Satisfaction: Enhanced service for top 91 customers
- 

## 1.9 8. Summary and Conclusions

### 1.9.1 Key Findings:

1. **Clustering Analysis:** Successfully identified [X] distinct customer segments with meaningful behavioral differences
2. **Dimensionality Reduction:** Visualizations confirmed cluster separation and revealed [linear/non-linear] data structure
3. **Anomaly Detection:** Identified [X] unusual customers requiring special attention, primarily

VIPs and high-risk cases

4. **Integration:** The three techniques complement each other, with dimensionality reduction confirming cluster validity and anomaly detection highlighting edge cases within clusters

### 1.9.2 Business Impact:

- **Targeted Marketing:** Each segment can receive customized campaigns
- **Risk Management:** Early identification of high-risk behaviors
- **Revenue Optimization:** VIP customers can receive premium services
- **Customer Retention:** Re-engagement strategies for inactive segments

### 1.9.3 Next Steps:

1. Implement segment-specific marketing campaigns
2. Monitor anomaly customers for fraud/opportunity
3. Track segment migration over time
4. A/B test recommendations to measure impact

```
[167]: # Save final results
df.to_csv('../results_with_clusters_and_anomalies.csv', index=False)
print("Analysis complete! Results saved to 'results_with_clusters_and_anomalies.
      csv'")
```

Analysis complete! Results saved to 'results\_with\_clusters\_and\_anomalies.csv'